



FABRIC UPGRADE PROCEDURE TO V8.3

This document describes:

- How to upgrade Fabric from the latest **V8.2** to the present **V8.3** version.
- How to re-implement the modified product features.

Notes:

- This document does not cover the Fabric server topology changes, such as nodes addition, DCs, change of replication factor and consistency level.
- It is a must to perform the Fabric upgrade procedure in the testing environments prior to applying it on your production deployment.
- Sanity tests must be performed upon the completion of the upgrade procedure, such as performing a few GET commands and conducting other checks per the sanity procedure defined in your project.

SOFTWARE UPGRADE PROCEDURE

Preliminary Step

Download the latest Fabric package release and copy it to the server(s).

Stop Fabric

Take the following steps in the specified order:

1. If your project has an iidFinder:

- Stop the iidFinder on all nodes.
- Wait until Kafka lags are zero in the relevant Kafka topics:
 - Delta_cluster_<LU_name> topics (e.g. deltaGroup)
 - DeltaPriority_cluster_<LU_name> topics (e.g. IDfinderGroupId_Cluster)
- Run the following command to verify that the lag on the above topics is zero:

```
/opt/apps/k2view/apps/kafka/bin/kafka-consumer-groups --bootstrap-server <internal Kafka server IP>
--group <Kafka interface group ID> --describe
```
- Investigate the remaining messages in the Delta tables and clean them, by taking the following steps per each LU:
 - Run MIGRATE command on all distinct IIDs.
 - Check the results in order to decide how to proceed with the failed entity messages.
 - Clean the Delta table.

2. Stop Fabric on all nodes.

Open the Package

Perform the following steps:

- Rename the Fabric directory as shown, type the specific Fabric version in the indicated location:

```
cp -r config config_${k2fabric -version | awk '{print $2}' | head -n1}
mv fabric ${k2fabric -version | awk '{print $2}' | head -n1}
```



FABRIC UPGRADE PROCEDURE

- Extract (un-TAR) the Fabric directories from the upgrade package (extract only the directories) as shown. The specific file name will depend on the specific Fabric version:

```
tar -zxvf k2fabric-server-fabric-<package_name>.tar.gz fabric
```

Verify Upgrade Success

Use the following command to verify that the upgrade procedure has been completed successfully:

```
k2fabric -version
```

The result will display the Fabric package number, for example:

```
Tag fabric- 8.2.0_220 at revision = 4c5e23ec2c464daea283faed9bceec2468d21225
```

If there are local files on the server (such as local JARs), their names will be displayed here.

Configuration Changes

Prior to performing the changes, backup your project configuration files, such as config.ini, etc. Please read the configuration changes summary per each version.

N/A

Implementation Changes

Please see the implementation changes summary for each version.

Upgrade To Version 8.3.0

Fabric Catalog

- Splitting the Catalog artifacts is now enabled by default due to the setting SPLIT_CATALOG_ARTIFACTS = ON in the config.ini file. As a result, separate artifacts are created per each data platform and schema, with the following name format: `catalog_field_info__<dataPlatform>_<schema>.csv`. The content of these files is automatically combined into one unified MTable `catalog_field_info` in Fabric's memory.
- Note that the existing `catalog_field_info.csv` file (that exists in your project) **needs to be deleted from the project**, to prevent data duplication in the combined MTable.

IID Finder

- If the project has customization for IID Finder, aka IIDF hook libraries, make sure these libraries are applied for the new fabric version – to `fabric/lib/provided/` location.

Start Fabric

Take the following steps in the specified order:

1. Start the first Fabric node and wait until it has been successfully completed.
2. Start all other nodes.
3. Deploy the project.
4. If your project has an iidFinder, re-start the iidFinder on all nodes.



FABRIC UPGRADE PROCEDURE

Note that the above steps may vary per your project's runbook. For example, you may first restart the iidFinder on a single node only, verify that the process has been successfully completed and then proceed to restarting the iidFinder on all other nodes.