



FABRIC V8.5.0 RELEASE NOTES

These Release Notes describe the new features in Fabric release V8.5.0 and list bugs that have been fixed since the latest V8.4 release.

Certification of this Fabric release is based on:

3 rd Party Name	Included with Fabric	3 rd Party Server
Cassandra	java-driver-core-4.17.0.jar	4.1.8
Elasticsearch	8.13.1	8.5.3
Kafka	3.9.1	confluent-8.1.1
Neo4j Enterprise	5.26.28	-
OpenJDK Runtime Environment	openjdk-21.0.11+10	openjdk-21.0.11+10
OpenSearch	3.4.0	AWS 1.3.4
PostgreSQL	postgresql-42.7.8.jar	17.5
SQLite	3.46.1	-

MAIN FEATURES AND IMPROVEMENTS

1. Fabric Catalog

The Discovery & Catalog solution has been enhanced with the following new features:

- **Large Catalog navigation:**
 - Navigating large schemas in the Catalog App could be slow and unresponsive. Expanding a schema with thousands of datasets, fields, and relations requires retrieving and rendering a large volume of data on the client side, which in some cases caused the browser to crash.
 - Introducing the **Schema Dataset List view**: for large schemas, datasets are displayed in a filterable list instead of the tree view, reducing rendering time and improving navigation.
 - All actions available in the standard tree view — such as exploring fields, viewing and editing dataset and field properties, viewing and adding relations — are fully supported in the list view as well.

https://support.k2view.com/Academy/articles/39_fabric_catalog/catalog_app/18_large_schema_navigation.html

- **File Cataloging — OpenAPI Support:**
 - The File Cataloging framework has been extended to support discovery of OpenAPI interfaces (versions 3.0 and 3.1).



FABRIC RELEASE NOTES

- A new **OpenApiToMetadata** Broadway actor is added to parse OpenAPI specifications and transform them into Catalog's standard metadata format, enabling the standard Discovery pipeline — including classification, PII detection, and more — to run in the same way as for any other data source.
- API endpoints (paths) are mapped as Catalog datasets, with schemas derived from OpenAPI tags and field sets merged across all HTTP methods defined on each path.

https://support.k2view.com/Academy/articles/39_fabric_catalog/06_open_api_support.html

- **Search Catalog enhancement:**

- The Catalog Advanced search has been enhanced with new 'Limit by hierarchy' search criteria, which allows users to limit the search scope by selecting a Data Platform or both a Data Platform and a Schema.

https://support.k2view.com/Academy/articles/39_fabric_catalog/catalog_app/08_search_catalog.html

- **Data Quality Metrics plugin:**

- Enhanced with additional metrics calculation and input parameters.
- Added a restriction related to PII fields: when a field is marked as PII, the minimum and maximum values are not calculated for it to avoid revealing sensitive data.

https://support.k2view.com/Academy/articles/39_fabric_catalog/plugins/04_source_data_metrics.html

- **Reference by Name Comparison plugin:**

- Introducing a new rule - *fieldNameCompare* - based on the Jaro-Winkler similarity algorithm, which compares two field name strings and produces a score between 0 (completely different) and 1 (identical). Unlike the other matching rules, this rule does not perform a PK check of the compared fields and can create a *refersTo* relation between any two fields based on name similarity alone.

https://support.k2view.com/Academy/articles/39_fabric_catalog/plugins/05_relations_creation.html

- **Reference by LLM plugin:**

- This new plugin identifies possible FK references between datasets by analyzing dataset pairs using an LLM and creates corresponding *refersTo* relations.

https://support.k2view.com/Academy/articles/39_fabric_catalog/plugins/05_relations_creation.html

- **Class nodes - naming convention improvement:**

- Until now, when an ad-hoc class was created in the Catalog — either from a complex field parsed from data (e.g., JSON, XML) or from a JSON Schema ad-hoc class definition — the parent class names used to be concatenated to the class name to prevent name ambiguity, e.g.:
 - AddressClass;PrimaryContactClass;CityClass
 - AddressClass;ShippingInfoClass;CityClass



FABRIC RELEASE NOTES

- In this version, this logic was changed to align with the Web Studio naming convention used during the LU Schema creation. The class name is now kept as is, and uniqueness is achieved by appending 3 underscores followed by the parent class name. If the parent's name alone doesn't provide uniqueness, additional ancestor levels (grandparent and beyond) are appended, separated by a single underscore, e.g.:
 - CityClass__AddressClass_PrimaryContactClass
 - CityClass__AddressClass_ShippingInfoClass
- Note that all classes sharing the same name receive the same number of appended ancestor names. When the root is reached, the @root@ annotation is used as the parent's name.
- **DISCOVERY_JDBC_LIMIT_BY_QUERY parameter refactoring:**
 - The DISCOVERY_JDBC_LIMIT_BY_QUERY parameter has been removed from the config.ini, while being replaced by another parameter – **Limit** – available in the Discovery Pipeline settings screen, under the Crawler section (default: false).
 - When Limit is set to true, the query limit is applied directly in the SQL query rather than via the JDBC driver's setMaxRows method.
- **Catalog Settings enhancements:**
 - The **Classifier Regex**, **PII & Masking**, and **Sequences tabs** in Catalog > Settings now offer a smoother filtering experience: the Classification filter is always expanded and focused, so a user can simply start typing to filter the list, without first having to click and select the filter field.
 - The settings of **Classifier Regex** and **PII & Masking** have been improved, by removing the regexes which caused high number of false positive, adding some more specific regexes (e.g. CREDIT_CARD, STATE), adding default generators for most classifications, etc.
 - The **Classifier Regex** tab allows now to disable a classification rule instead of deleting it, making it easier to manage rules that are temporarily unused. The disabled rules are skipped by the *Metadata Regex Classifier* and *Data Regex Classifier* classification plugins.

https://support.k2view.com/Academy/articles/39_fabric_catalog/catalog_app/10_catalog_settings.html

2. K2exchange:

- OpenAI / Anthropic connectors V2.0 were released and should be re-installed in the existing projects.

3. Security Improvements

- **API Key security:** Strengthen API Key-based authentication by giving a stronger, higher-entropy key option. API Keys now support three explicit types, selectable via CREATE TOKEN or the Web Studio Admin UI:



FABRIC RELEASE NOTES

- Access - high-entropy key, when sending API key as *Authorization: Bearer* header. This is the recommended type and is now the default in the Admin UI.
- Legacy - the existing name-as-key behavior, preserved for backward compatibility and thus also set as default.
- Signing JWT - the client signs its own short-lived JWT using a Fabric-issued signing key (replacing the previous "Signed by client" option).

Both the Access and Signing JWT keys are never displayed or retrievable again after creation.

- **Authenticate API mechanism** - The /authenticate endpoint now distinguishes end-user logins from server-to-server API calls. End-user sessions continue to receive their JWT as a secure cookie tied to their identity. Server-to-server callers now explicitly request the JWT in the response body, keeping API access decoupled from browser session cookies. Audit logs now clearly attribute each login to the resolved user or, for pure API-key access, to the API Key itself.

4. Other Improvements and Changes

- The SET INSTANCE_TTL command has been enhanced to support Time To Live (TTL) management for MicroDBs stored in **Google Cloud Storage (GCS)**, **Azure Blob Storage**, and **Amazon S3**.

https://support.k2view.com/Academy/articles/02_fabric_architecture/04_fabric_commands.html

- **Hybrid microDB cache**

- A new CACHE_TYPE value, HYBRID_NO_CACHE, is available. This option is designed for data products that contain a mix of small and large MicroDBs, optimizing both performance and memory utilization.
- When CACHE_TYPE is set to HYBRID_NO_CACHE, Fabric determines where to store each MicroDB cache, in memory or disk, based on its size. The size threshold is configured using the HYBRID_MEMORY_THRESHOLD parameter in the [fabricdb] section.

https://support.k2view.com/Academy/articles/32_LU_storage/05_in_memory_lu_storage.html

- **Affinity Management:**

- Fabric now supports centralized management of node affinities across the cluster. Instead of editing the node.id file on each node separately, affinity rules can be defined once - per DC or for the entire cluster - via the Admin UI or the new **set_global affinity_rules** command.
- A dedicated Affinity Management Job (one per DC) automatically distributes and maintains the settings on all nodes.
- The feature also introduces node groups (colors) and new commands for setting, viewing and monitoring node affinities.

https://support.k2view.com/Academy/articles/20_jobs_and_batch_services/19_affinity_management.html



FABRIC RELEASE NOTES

- **Retry during parallel download from S3 / Azure**
 - If during the parallel download from S3 or Azure the BLOB is modified and saved, the retry is triggered automatically to download the updated BLOB.
- **Support of S3 versioning during parallel download**
 - New parameter `#IGNORE_VERSION_CHANGE_DURING_PARALLEL_DOWNLOAD` has been added to the `config.ini` and by default it's set to `true`.
 - When versioning is enabled on the S3 bucket, this parameter can be set to **false**, to enable parallel download of a specific BLOB version.
- **Batch size**
 - A new parameter `BATCH_FETCH_SIZE=1000` has been added to the `[batch_process]` section of the `config.ini`. This parameter defines the fetch size which should be used when batch population is based on SQL query. Its purpose is to improve the performance of the IID generation step of the Batch process.
- **Wrapping fabric generated queries with quotes**
 - When the parameter **`ALWAYS_QUOTES_INTERNAL_QUERIES`** is set to `true`, all fields in fabric generated queries are wrapped with quotes.
 - Now the default is set to **false**.
- **Enhanced the DROP LUTYPE Fabric command to support dropping only the data.** A new `data_only` parameter has been added; set it to `Y` to drop only the data while preserving the LUTYPE definition.
- **List of schemas per interface**
 - The existing `/api/interface_schema/schemas` REST API that retrieves a list of interface schemas has been enhanced to return all interfaces with an indication whether it's retrieved from Catalog or not. For that purpose, a new `schemaSource` input param has been added and it can get one of the following values:
 - `<empty>` (default) - if the interface is in the Catalog, bring schemas from Catalog; otherwise - bring them from the source. The same behavior as before V8.5.
 - **source** - bring schemas from source only.
 - **catalog** - bring schemas from Catalog only.
 - **both** - bring schemas from Catalog and from the source.
- From now on, there is no need to restart Fabric when setting `AUDIT=ON` in the `config.ini`.
- **Improved MDB export performance** by using the PostgreSQL `COPY` command when exporting LU instances to PostgreSQL.



FABRIC RELEASE NOTES

- Enhanced the Web Server shutdown process to support graceful termination of active requests. During shutdown, the server stops accepting new connections, allowing load balancers to redirect incoming traffic to other nodes, while continuing to process existing requests for a configurable grace period. The grace period is controlled by the `WEB_SERVICE_SHUTDOWN_TIMEOUT_SECONDS` parameter (default: **30 seconds**). Once all active requests are completed, shutdown proceeds immediately without waiting for the full timeout. Requests that exceed the configured grace period are terminated, and any unfinished request-processing threads are logged with warning messages and stack traces to assist with troubleshooting. Setting the timeout to **0** disables the grace period and performs an immediate shutdown. This enhancement minimizes customer impact during planned restarts and improves high-availability behavior in load-balanced environments.

RESOLVED ISSUES

- The discovery LLM Description plugin was triggering a data snapshot step, even when its sample size was set to 0 and all other data-related plugins were inactive.
- Ticket #45272 – JobReconcile got stuck on quitting and didn't execute any more jobs. It was caused by InnerFlowAsync actor in one of the Broadway jobs.
- Ticket #48292 – Broadway iterated over result set of disabled DbCommand actor.
- Ticket #49345 – In the Admin web app, cluster time showed UTC time zone instead of the local one.
- Ticket #49704 – In Broadway, creating links between the actors by using the link manager didn't work.
- Ticket #50261 – The Web Framework Admin menu remained visible despite `hidden=true` in `apps.json`. Now `apps.json` properties (like "hidden"), defined and maintained at the project, overrides the base `apps.json` properties.
- Ticket #50447 – Catalog parser didn't recognize CLOB field in Oracle DB with XML inside as complex.
- Ticket #50721 – when running Discovery for a table with a CLOB field with multiple JSON structures and one of the JSON's inner levels sometimes includes an object and sometimes is null, this object was not parsed as complex.
- Following 3rd party tools were upgraded:
 - Tomcat 11.0.22
 - Neo4j 5.26.28
 - MCP core 1.1.1
 - Masterkey Fortanix log4j-core-2.25.4
 - rhino-1.7.7.2.jar – removed
 - google-oauth-client-1.39.0